

compareMCMCs: An R package for studying MCMC efficiency

Perry de Valpine¹, Sally Paganin¹, and Daniel Turek²

¹ University of California, Berkeley ² Williams College

DOI: [10.21105/joss.03844](https://doi.org/10.21105/joss.03844)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: Fabian Scheipl ↗

Reviewers:

- [@rowlandseymour](#)
- [@tbrown122387](#)

Submitted: 18 October 2021

Published: 22 October 2021

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Markov chain Monte Carlo (MCMC) algorithms are used to stochastically simulate from complicated probability distributions. MCMC is very widely used to implement Bayesian statistical analysis, where the distribution of interest (the “target distribution”) is a posterior distribution of parameters given data. In this context, the posterior is known only up to a constant, so only relative probabilities (or densities) can be easily calculated, which is sufficient for MCMC to work. In applied Bayesian statistics, MCMC algorithms for large or complex statistical models and data sets are sometimes run for minutes, hours or days, making them an analysis bottleneck, so there is a premium on efficiency. MCMC efficiency includes both computational speed and algorithmic mixing, which refers to how well the algorithm explores the posterior distribution from one iteration to the next. Computational speed may comprise one or more steps such as algorithm setup, MCMC “burn-in” or “warm-up” phases, and MCMC execution or “sampling.”

There are many MCMC algorithms (also called “samplers”) and software packages implementing them. Because MCMC samplers can be mixed and matched, with different samplers for the same or different dimensions of a target distribution possibly combined into the same MCMC algorithm, there is an enormous space of MCMC algorithms. Invention of new methods, comparisons among methods, and theoretical study of MCMC mixing are all important areas of active research. Various software packages provide samplers such as Gibbs, adaptive random-walk Metropolis-Hastings, slice, Hamiltonian, multivariate (“block”) or other variants of these, and others. Different MCMC algorithms can yield efficiency that differs by orders of magnitude for a particular problem, with these variations in efficiency being problem-dependent.

The R package `compareMCMCs` provides a highly modular system for managing performance comparisons among MCMC software packages for purposes of research on MCMC methods. MCMC runs can take a long time, so the output (“samples”) and components of computation time from a run are stored regardless of whether performance metrics are computed immediately. Arbitrary MCMC packages (“MCMC engines”) can be added to the system by writing a simple plug-in or wrapper to manage inputs and outputs in a unified way. Conversions among model parameter names and/or different parameterizations can be provided to standardize across packages. Performance metrics are organized by model parameter (one result per parameter per MCMC engine), by MCMC (one result per MCMC engine), or arbitrarily (a user-defined list of metric results per MCMC engine). Built-in metrics include two methods of estimating effective sample size (ESS), posterior summaries such as mean and common quantiles, efficiency defined as ESS per computation time, rate defined as computation time per ESS, and minimum efficiency per MCMC. New metrics can be provided by a plug-in system and applied programmatically to a set of MCMC samples without re-running the MCMC engines. Finally, standardized graphical comparison pages can be generated in html. Built-in graphical outputs include figures comparing MCMC efficiency and/or rate on a per-parameter or per-MCMC basis as well as comparing posterior distributions. New graphical

44 outputs can be provided by a plug-in system. In summary, `compareMCMCs` is modular and
45 extensible for running new MCMC engines on comparable problems, for creating new metrics
46 of interest (e.g., posterior summaries or effective sample size estimated in different ways), and
47 for creating new graphical comparison outputs into a report.

48 Use of `compareMCMCs` supports but does not require a primary role for MCMCs created with
49 the `nimble` (de Valpine et al., 2017, 2021) package for hierarchical statistical models. That
50 is because `nimble` provides greater flexibility than other packages to customize its MCMC
51 system, configuring which samplers will operate on which parts of a model and/or writing
52 new samplers. Thus, it is of interest to compare multiple MCMC methods all implemented
53 within `nimble`. Furthermore, `nimble` uses a model language that is a dialect of that used
54 by WinBUGS, OpenBUGS, MultiBUGS, and JAGS (Goudie et al., 2020; D. Lunn et al., 2009;
55 D. J. Lunn et al., 2000; Plummer & others, 2003). These packages are often called from R
56 via packages such as R2WinBUGS (Sturtz et al., 2005), `rjags` (Plummer, 2019), and `jagsUI`
57 (Kellner, 2019). Therefore, for fully compatible models, comparisons between `nimble` and
58 JAGS can be run in `compareMCMCs` from the same model and data specifications. A plug-in is
59 also provided for Stan via `rstan` (Stan Development Team, 2020), and the extension system
60 to plug in new MCMC engines is clearly documented.

61 Statement of need

62 Many other packages run MCMC algorithms and/or post-process MCMC results, but `compa`
63 `reMCMCs` is distinct in its goal of supporting MCMC research by comparing MCMC methods.
64 Packages that run MCMC from R are documented on the “Cran Task View” page for “Bayesian
65 Inference” (Park, 2021) of the Comprehensive R Archive Network (CRAN). Some popular
66 general packages include those listed above as well as others such as `MCMCpack` (Martin et
67 al., 2011) and `LaplacesDemon` (Statisticat & LLC., 2021). Furthermore, there are MCMC
68 engines based in Python, such as `PyMC` (Salvatier et al., 2016), and other languages. These
69 may be called via appropriate interfaces from R to other languages.

70 Of the packages listed on the “Bayesian Inference” Task View, only the `SamplerCompare`
71 (Thompson, 2011) package appears to specifically support the goal of comparing MCMC
72 performance. However, this package can only compare MCMC samplers that have exactly
73 one scalar tuning parameter, target distributions that are continuous with constant dimension,
74 and are implemented within the package.

75 Packages for post-processing of MCMC samples (e.g., `coda` (Plummer et al., 2006), `BayesP`
76 `ostEst` (Scogin et al., 2019), and `MCMCvis` (Youngflesh, 2018)) aim to provide features for
77 scientific summary and presentation of results, whereas `compareMCMCs` provides features for
78 comparisons of algorithm performance across packages. Assessing MCMC performance is not
79 simply a matter of computational benchmarking. For example, effective sample size is itself a
80 non-trivial property to estimate by statistical methods, different metrics may be of interest for
81 different purposes, and consistency of algorithm results between different MCMC engines can
82 only be determined statistically, i.e. within simulation error. Therefore, the features needed for
83 comparing MCMC performance are distinct from those needed for presenting scientific results
84 based on MCMC.

85 Acknowledgements

86 We thank other developers who have contributed to `nimble`, in which an early version of the
87 features of `compareMCMCs` was first drafted and from which its design needs became clearer.
88 Daniel Turek wrote the first version of what evolved into `compareMCMCs`, and Christopher
89 Paciorek contributed ideas along the way.

References

- de Valpine, P., Paciorek, C., Turek, D., Michaud, N., Anderson-Bergman, C., Obermeyer, F., Wehrhahn Cortes, C., Rodríguez, A., Temple Lang, D., & Paganin, S. (2021). *NIMBLE: MCMC, particle filtering, and programmable hierarchical modeling* (Version 0.11.1) [Computer software]. <https://doi.org/10.5281/zenodo.1211190>
- de Valpine, P., Turek, D., Paciorek, C., Anderson-Bergman, C., Temple Lang, D., & Bodik, R. (2017). Programming with models: Writing statistical algorithms for general model structures with NIMBLE. *Journal of Computational and Graphical Statistics*, 26, 403–413. <https://doi.org/10.1080/10618600.2016.1172487>
- Goudie, R. J. B., Turner, R. M., De Angelis, D., & Thomas, A. (2020). MultiBUGS: A parallel implementation of the BUGS modelling framework for faster bayesian inference. *Journal of Statistical Software*, 95(7). <https://doi.org/10.18637/jss.v095.i07>
- Kellner, K. (2019). *jagsUI: A wrapper around 'rjags' to streamline 'JAGS' analyses*. CRAN. <https://cran.r-project.org/package=jagsUI>
- Lunn, D. J., Thomas, A., Best, N., & Spiegelhalter, D. (2000). WinBUGS—a bayesian modelling framework: Concepts, structure, and extensibility. *Statistics and Computing*, 10(4), 325–337.
- Lunn, D., Spiegelhalter, D., Thomas, A., & Best, N. (2009). The BUGS project: Evolution, critique and future directions. *Statistics in Medicine*, 28(25), 3049–3067.
- Martin, A. D., Quinn, K. M., & Park, J. H. (2011). MCMCpack: Markov chain monte carlo in R. *Journal of Statistical Software*, 42(9), 22. <https://www.jstatsoft.org/v42/i09/>
- Park, J. H. (2021). *CRAN Task View: Bayesian Inference*. CRAN. <https://CRAN.R-project.org/view=Bayesian>
- Plummer, M. (2019). *Rjags: Bayesian graphical models using MCMC*. CRAN. <https://cran.r-project.org/package=rjags>
- Plummer, M., Best, N., Cowles, K., & Vines, K. (2006). CODA: Convergence diagnosis and output analysis for MCMC. *R News*, 6(1), 7–11. <https://journal.r-project.org/archive/>
- Plummer, M., & others. (2003). JAGS: A program for analysis of bayesian graphical models using gibbs sampling. *Proceedings of the 3rd International Workshop on Distributed Statistical Computing*, 124, 1–10.
- Salvatier, J., Wiecki, T. V., & Fonnesbeck, C. (2016). Probabilistic programming in python using PyMC3. *PeerJ Computer Science*, 2, e55. <https://doi.org/10.7717/peerj-cs.55>
- Scogin, S., Karreth, J., Beger, A., & Williams, R. (2019). BayesPostEst: An r package to generate postestimation quantities for bayesian MCMC estimation. *Journal of Open Source Software*, 4(42), 1722. <https://doi.org/10.21105/joss.01722>
- Stan Development Team. (2020). *RStan: The R interface to Stan*. <http://mc-stan.org/>
- Statisticat, & LLC. (2021). *LaplacesDemon: Complete environment for bayesian inference*. Bayesian-Inference.com. <https://web.archive.org/web/20150206004624/http://www.bayesian-inference.com/software>
- Sturtz, S., Ligges, U., & Gelman, A. (2005). R2WinBUGS: A package for running WinBUGS from r. *Journal of Statistical Software*, 12(3), 1–16. <http://www.jstatsoft.org>
- Thompson, M. B. (2011). Introduction to SamplerCompare. *Journal of Statistical Software*, 43(12), 1–10. <http://www.jstatsoft.org/v43/i12/>
- Youngflesh, C. (2018). MCMCvis: Tools to visualize, manipulate, and summarize MCMC output. *Journal of Open Source Software*, 3(24), 640. <https://doi.org/10.21105/joss.00640>