

Castor: A C++ library to code “à la matlab”

Matthieu Aussenat^{*1}, Marc Bakry^{†1}, and Laurent Series^{‡2}

¹ Ecole Polytechnique (CMAP), INRIA, Institut Polytechnique Paris, Route de Saclay 91128, Palaiseau, France ² Ecole Polytechnique (CMAP), Institut Polytechnique Paris, Route de Saclay 91128, Palaiseau, France

DOI: [10.21105/joss.03965](https://doi.org/10.21105/joss.03965)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Mehmet Hakan Satman](#)

↗

Reviewers:

- [@mkitti](#)
- [@pitsianis](#)

Submitted: 23 November 2021

Published: 11 December 2021

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

The objective of the *Castor* framework is to propose high-level semantics, inspired by the Matlab language, allowing fast software prototyping in a low-level compiled language. It is nothing more than a matrix management layer using the tools of the standard C++ library (C++14 and later), in different storage formats (full, sparse and hierarchical). Linear algebra operations are built over the BLAS API and graphic rendering is performed in the VTK framework. The *Castor* framework is provided as an open source software under the LGPL 3.0.

Statement of need

Matlab is a software used worldwide in numerical prototyping, due to its particularly user-friendly semantics and its certified toolboxes. However, some usecases do not allow codes in Matlab format, for example multi-platform portability issues or propriety licensing. There is many open source libraries dealing with dense / sparse matrices, but none to our knowledge offers the same semantic as Matlab for manipulations and also both for algebra and visualization. To start meeting these needs, a header-only template library for matrix management has been developed, based on the standard C++ library, by encapsulating the `std::vector` class. Many tools and algorithms are provided to simplify the development of scientific computing programs:

- dense, sparse and hierarchical matrices manipulations,
- linear algebra computations (optimized BLAS library),
- graphical representations (VTK library).

Compared to standard C++, this high-level semantic/low-level language coupling makes it possible to gain efficiency in the prototyping phase, while ensuring performance for applications. In addition, direct access to data structures allows users to optimize the most critical parts of their code. Finally, a complete documentation is available, as well as continuous integration unit tests. All of this makes it possible to meet the needs of teaching (notebooks using c++ interpreter such as Cling), academic research and industrial applications at the same time.

^{*}matthieu.aussenat@polytechnique.edu

[†]marc.bakry@polytechnique.edu

[‡]laurent.series@polytechnique.edu

34 Dense Matrix

35 The dense matrix part of the *Castor* framework implements its own templated class `matrix<T>` in `matrix.hpp`, where `T` can be fundamental types of C++ as well as `std::complex`.
 36 This class is built over a `std::vector<T>` which holds the values (ISO/IEC, 2014). Note
 37 that the element of a matrix is stored in row-major order and that a vector is considered as a
 38 $1 \times n$ or $n \times 1$ matrix.

40 The class `matrix<T>` provides many useful functions and operators such as:

- 41 ■ builders which can be used to initialize all coefficients (zeros, ones, eye, etc.),
- 42 ■ standard algorithms over data stored in matrices (norm, max, sort, argsort, etc.),
- 43 ■ mathematical functions which can be applied element-wise (cos, sqrt, conj, etc.),
- 44 ■ matrix manipulations like concatenate matrices in all dimensions, find the non-zero
 45 elements or transpose them, reshape size, etc.,
- 46 ■ standard C++ operators which have been overloaded and work element-wise (+, *, !, &
 47 &etc.),
- 48 ■ values accessors and matrix views with linear and bi-linear indexing,
- 49 ■ elements of linear algebra, such as the matrix product (mtimes or tgemm) and linear
 50 system resolution (multi-right-hand-side gmres),
- 51 ■ many other tools to display elements (<<, disp), save and load elements from file
 52 (ASCII or binary), etc.

53 The API provides more than a hundred functions and is designed such that it should feel like
 54 using Matlab. For advanced users, direct access to the data stored in the `std::vector<T>`
 55 enables all or part of an algorithm to be optimized in native C++.

56 This example displays the sum of two matrices with implicit cast :

```
#include <iostream>
#include "castor/matrix.hpp"
using namespace castor;
int main(int argc, char* argv[])
{
    matrix<float> A = {{ 1.0, 2.0, 3.0, 4.0},
                      { 5.0, 6.0, 7.0, 8.0},
                      { 9.0, 10.0, 11.0, 12.0}};
    matrix<double> B = eye(3,4);
    auto C = A + B;
    disp(C);
    return 0;
}
```

```
57 Matrix 3x4 of type 'd' (96 B):
58      2.00000      2.00000      3.00000      4.00000
59      5.00000      7.00000      7.00000      8.00000
60      9.00000     10.00000     12.00000     12.00000
```

61 Linear Algebra

62 The linear algebra part of the framework, implemented in `linalg.hpp`, provides a set of useful
 63 functions to perform linear algebra computations by linking to optimized implementations of
 64 the BLAS and LAPACK standards (Anderson et al., 1999) (OpenBLAS, oneAPI MKL, etc.).

65 The BLAS part is a straightforward overlay of the C-BLAS type III API, which is compatible
66 with row-major ordering. This is achieved by a template specialization of the `tgemm` function,
67 which allows optimized implementations to take control of the computation using `sgemm`, `dgem`
68 `m`, `cgemm` and `zgemm`. Thanks to this interface, naive implementations proposed in *matrix.hpp*
69 for dense matrix-products `mtimes` and `tgemm` may be improved in term of performance,
70 especially for large matrices.

71 The LAPACK part is a direct overlay over the Fortran LAPACK API, which uses a column
72 ordering storage convention. This interface brings new high-level functionalities, such as a
73 linear solver (`linsolve`), matrix inversion (`inv`, `pinv`), factorizations (`qr`, `lu`), the search
74 for eigen or singular values decompositions (`eig`, `svd`), `aca` compression (`aca`), etc. It uses
75 templated low-level functions following the naming convention close to the LAPACK one
76 (like `tgesdd`, `tgeqrf`, etc.).

77 This example displays the product of A and A^{-1} :

```
#include <iostream>
#include "castor/matrix.hpp"
#include "castor/linsolve.hpp"
using namespace castor;
int main (int argc, char* argv[])
{
    matrix<> A = rand(4);
    matrix<> Am1 = inv(A);
    disp(mtimes(A,Am1));
    return 0;
}
```

```
78 Matrix 4x4 of type 'd' (128 B):
79 1.0000e+00  1.0408e-16 -2.7756e-17 -5.5511e-17
80      0  1.0000e+00 -5.5511e-17  1.1102e-16
81      0 -2.2204e-16  1.0000e+00 -1.1102e-16
82 -2.7756e-17      0      0  1.0000e+00
```

83 2D/3D Visualization

84 The graphic rendering part, provided by *graphics.hpp*, features 2D/3D customizable plotting
85 and basic mesh generation. It is based on the well-known VTK library (Schroeder et al.,
86 2000). Here again, the approach tries to get as close as possible to Matlab semantics.

87 First, the user creates a `figure`, which is a dynamic container of data to display. The `figure`
88 class is composed of a `vtkContextView` class, providing a view with a default interactor
89 style, renderer, etc. Then, graphic representations can be added to the figure, using functions
90 like `plot`, `imagesc`, `plot3`, `mesh`, etc. Options are available to customize the display of
91 the results, such as the plotting style, legend, colorbar and others basic stuff. Finally, the
92 `drawnow` function must be called to display all defined figures. The latters are displayed and
93 manipulated in independent windows.

94 In addition, graphics exports are available in different compression formats (`png`, `jpg`, `tiff`,
95 etc.), as well as video rendering (`ogg`).

96 This example shows a basic 2D plotting of a sine function (Figure 1):

```
#include "castor/matrix.hpp"
#include "castor/graphics.hpp"
```

```
using namespace castor;
int main (int argc, char* argv[])
{
    matrix<> X = linspace(0,10,100);
    figure fig;
    plot(fig,X,sin(X),{"r-+"},{"sin(x)"});
    plot(fig,X,cos(X),{"bx"},{"cos(x)"});
    drawnow(fig);
    return 0;
}
```




Figure 1: Basic 2D plotting from Castor (using VTK).

Sparse matrices

Some matrices have sparse structures, with many (many) zeros that do not need to be stored (Davis & Hu, 2016). There are adapted storage formats for this type of structure (LIL, COO, CSR, etc.), the most natural being to store the indices of rows and columns for each non-zero value, as a list of triplet $\{i, j, v\}$. For the *Castor* framework, a dedicated template class to this kind of matrix has been developed (see *smatrix.hpp*). The storage format is based on a row major sorted linear indexing. Only non-zero values and their sorted linear indices are stored in a list of pairs $\{v, l\}$: for a $m \times n$ matrix, the following bijection is used to switch with the common bilinear indexation:

- $\{i, j\} \rightarrow l = i \cdot n + j,$
- $l \rightarrow \{i = \frac{l}{n}; j = l \bmod n\}.$

Accessors to all the elements are provided so that sparse matrices can be manipulated in a similar way as the dense matrices. This operation is performed by dichotomy with a convergence in $\log_2(\text{nnz})$, where nnz is the number of non-zero elements. Just like dense matrices, numerical values are stored in a templated `std::vector<T>`. For convenience, we provide classical

112 builders (sparse, speye, spdiags, etc.), standard C++ operators overloading, views, display
113 functions (disp, spy) and some linear algebra tools (transpose, mtimes, gmres, etc.).

114 This example displays the sum of two sparse matrices, with implicit cast and sparse to dense
115 conversion :

```
#include <iostream>
#include "castor/smatrix.hpp"
using namespace castor;
int main (int argc, char* argv[])
{
    smatrix<float> As = {{0.0,  0.0,  0.0},
                       {5.0,  0.0,  7.0}};

    As(0,1) = 2.0;
    smatrix<double> Bs = speye(2,3);
    disp(As);
    disp(As(0,1)); // bilinear accessor
    disp(As(4));   // linear accessor
    disp(Bs);
    disp(full(As+B));
    return 0;
}
```

116 Sparse matrix 2x3 of type 'f' with 3 elements (12 B):

```
117 (0,1)  2
118 (1,0)  5
119 (1,2)  7
120 2
121 0
```

122 Sparse matrix 2x3 of type 'd' with 2 elements (16 B):

```
123 (0,0)  1
124 (1,1)  1
```

125 Matrix 2x3 of type 'd' (48 B):

```
126 1.00000  2.00000  0
127 5.00000  1.00000  7.00000
```

128 Hierarchical matrices

129 To widen the field of applications, the \mathcal{H} -matrix format, so-called hierachical matrices (Hack-
130 busch, 1999), have been added in *hmatrix.hpp*. They are specially designed for matrices with
131 localized rank defaults. It allows a fully-populated matrix to be assembled and stored in a
132 lighter format by compressing some parts of the original dense matrix using a low-rank repre-
133 sentation (Rjasanow, 2002). They are constructed by binary tree subdivisions in a recursive
134 way, with a parallel assembly of the compressed and full leaves (using the OpenMP standard).
135 This format features a complete algebra, from elementary operations to matrix inversion. An
136 example is given in the application section that follows.

137 Application with a FEM/BEM simulation

138 As an application example, an acoustical scattering simulation was carried out using a bound-
139 ary element method (BEM) tool, implemented with the *Castor* framework (see the *fembem*

140 package (Aussal & Bakry, 2021)). We consider a smooth n -oriented surface Γ of some ob-
 141 ject Ω , illuminated by an incident plane wave u_i with wave-number k . The scattered field u
 142 satisfies the Helmholtz equation in Ω , Neumann boundary conditions (*sound-hard*) and the
 143 Sommerfeld radiation condition:

- 144 ▪ $-(\Delta u + k^2 u) = 0$,
- 145 ▪ $-\partial_n u_i = 0$,
- 146 ▪ $\lim_{r \rightarrow +\infty} r (\partial_r u - iku) = 0$.

147 The scattered field u satisfies the integral representation (Neumann interior extension, see
 148 (Terrasse & Abboud, 2013)):

$$u(\mathbf{x}) = - \int_{\Gamma} \partial_{n_y} G(\mathbf{x}, \mathbf{y}) \mu(\mathbf{y}) d_y \quad \forall \mathbf{x} \in \mathbb{R}^3 \setminus \overline{\Omega},$$

149 for some density μ , with the Green kernel $G(\mathbf{x}, \mathbf{y}) = \frac{e^{ik|\mathbf{x}-\mathbf{y}|}}{4\pi|\mathbf{x}-\mathbf{y}|}$. Using the boundary conditions
 150 we obtain :

$$-H\mu(\mathbf{x}) = -\partial_n u_i(\mathbf{x}) \quad \forall \mathbf{x} \in \Gamma,$$

151 where the hypersingular operator H is defined by:

$$H\mu(\mathbf{x}) = \int_{\Gamma} \partial_{n_x} \partial_{n_y} G(\mathbf{x}, \mathbf{y}) \mu(\mathbf{y}) d_y.$$

152 The operator H is assembled using a P_1 finite element discretization on a triangular mesh of
 153 the surface Γ , stored using dense matrices (*matrix.hpp*) or hierarchical matrices (*hmatrix.hpp*).

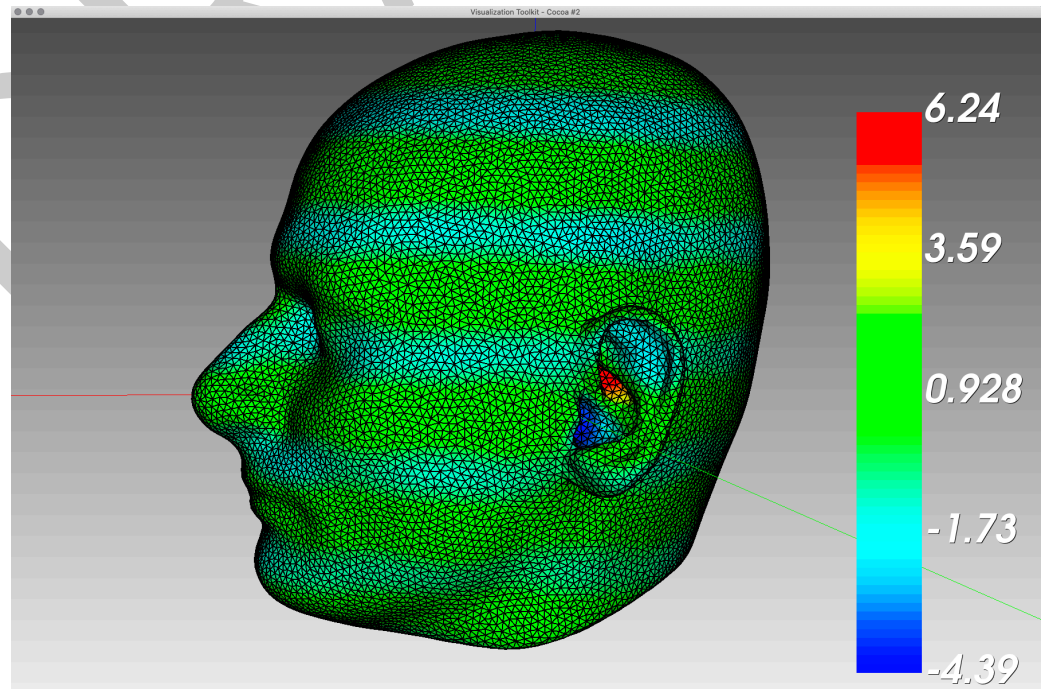


Figure 2: Resonance mode at 8kHz of the human pinna (BEM with H-Matrix).

154 Finally, using all the tools provided by Castor to write and solve these equations, we are able
155 to efficiently compute the acoustic diffraction of a harmonic plane wave at 8kHz, on a human
156 head mesh (Jin et al., 2013)). The simulation result (Figure 2) highlights the role of the
157 auditory pavilion as a resonator, modifying the timbre of a sound source to allow a listener's
158 brain to precisely locate its direction.

```
#include <castor/matrix.hpp>
#include <castor/smatrix.hpp>
#include <castor/hmatrix.hpp>
#include <castor/linalg.hpp>
#include <castor/graphics.hpp>
#include "fem.hpp"
#include "bem.hpp"

using namespace castor;

int main (int argc, char* argv[])
{
    // Load meshes
    matrix<double> Svtx;
    matrix<size_t> Stri;
    std::tie(Stri,Svtx) = triread("./","Head03_04kHz.ply");

    // Graphical representation
    figure fig;
    trimesh(fig,Stri,Svtx);

    // Parameters
    matrix<double> U = {0,0,-1};
    double f = 2000;
    double k = 2*M_PI*f/340;
    float tol = 1e-3;

    // FEM and mass matrix (sparse storage)
    tic();
    femdata<double> v(Stri,Svtx,lagrangeP1,3);
    femdata<double> u(Stri,Svtx,lagrangeP1,3);
    auto Id = mass<std::complex<double>>(v);
    toc();

    // Left hand side : [-H]
    tic();
    auto LHSfct = [&v,&u,&k](matrix<std::size_t> Ix, matrix<std::size_t> Iy)
    {
        return -hypersingular<std::complex<double>>(v,u,k,Ix,Iy);
    };
    hmatrix<std::complex<double>> LHS(v.dof(),u.dof(),tol,LHSfct);
    toc();
    disp(LHS);

    // Right hand side : - \int_{\Sigma_x} v(x) PW(x,u) dx
    auto B = - rightHandSide<std::complex<double>>(v,dnPWsource,U,k);

    // Solve -H = -dnP0
    hmatrix<std::complex<double>> Lh,Uh;
```



```

tic();
std::tie(Lh,Uh) = lu(LHS,1e-1);
toc();
disp(Lh);
disp(Uh);
auto mu = gmres(LHS,B,tol,100,Lh,Uh);

// Boundary radiation
tic();
auto Dbndfct = [&v,&u,&k,&Id](matrix<std::size_t> Ix, matrix<std::size_t> Iy)
{
    return 0.5*eval(Id(Ix,Iy)) + doubleLayer<std::complex<double>>(v,u,k,Ix,Iy);
};
hmatrix<std::complex<double>> Dbnd(v.dof(),u.dof(),tol,Dbndfct);
matrix<std::complex<double>> Pbnd = mtimes(Dbnd,mu);
toc();
Pbnd = - gmres(Id,Pbnd,tol,100) + planeWave(v.dof(),U,k);

// Graphical representation
figure fig2;
trimesh(fig2,Stri,Svtx,abs(Pbnd));

// Export in .vtk file
triwrite("./","head.vtk",Stri,Svtx,real(Pbnd));

// Plot
drawnow(fig);
disp("done !");
return 0;
}

```

159 Acknowledgements

160 We thank Houssem Haddar for his precious help.

161 References

- 162 Anderson, E., Bai, Z., Bischof, C., Blackford, S., Demmel, J., Dongarra, J., Du Croz, J.,
 163 Greenbaum, A., Hammarling, S., McKenney, A., & Sorensen, D. (1999). *LAPACK users'*
 164 *guide* (Third). Society for Industrial; Applied Mathematics. ISBN: 0-89871-447-8 ([paper-](#)
 165 [back](#))
- 166 Aussal, M., & Bakry, M. (2021). <https://gitlab.labos.polytechnique.fr/leprojetcastor/fembem>
- 167 Davis, T. A., & Hu, Y. (2016). The university of florida sparse matrix collection. *ACM*
 168 *Transactions on Mathematical Software (TOMS)*, 38(1), 1-25. [https://doi.org/10.1145/](https://doi.org/10.1145/2049662.2049663)
 169 [2049662.2049663](https://doi.org/10.1145/2049662.2049663)
- 170 Hackbusch, W. (1999). A sparse matrix arithmetic based on H-matrices. Part 1: Introduction
 171 to H-matrices. *Computing*, 62(2), 89–108.
- 172 ISO/IEC. (2014). International standard ISO/IEC 14882:2014(e) – programming language
 173 c++. Geneva, Switzerland: International Organization for Standardization.

- 174 Jin, C. T., Guillon, P., Epain, N., Zolfaghari, R., Van Schaik, A., Tew, A. I., & Thorpe, J.
175 (2013). Creating the sydney york morphological and acoustic recordings of ears database.
176 *IEEE Transactions on Multimedia*, 16(1), 37-46. <https://doi.org/10.1109/icme.2012.93>
- 177 Rjasanow, S. (2002). Adaptive cross approximation of dense matrices. *Int. Association*
178 *Boundary Element Methods Conf., IABEM* (pp. 28-30).
- 179 Schroeder, W. J., Avila, L. S., & Hoffman, W. (2000). Visualizing with VTK: A tutorial. *IEEE*
180 *Computer Graphics and Applications*, 20(5), 20-27. <https://doi.org/10.1109/38.865875>
- 181 Terrasse, I., & Abboud, T. (2013). Modélisation des phénomènes de propagation d'ondes.
182 *École Polytechnique*.

DRAFT